

23 - April

# A Deep Reinforcement Learning Framework For Identifying Funny Scenes In Movie

**DA 671** – Introduction to Reinforcement Learning

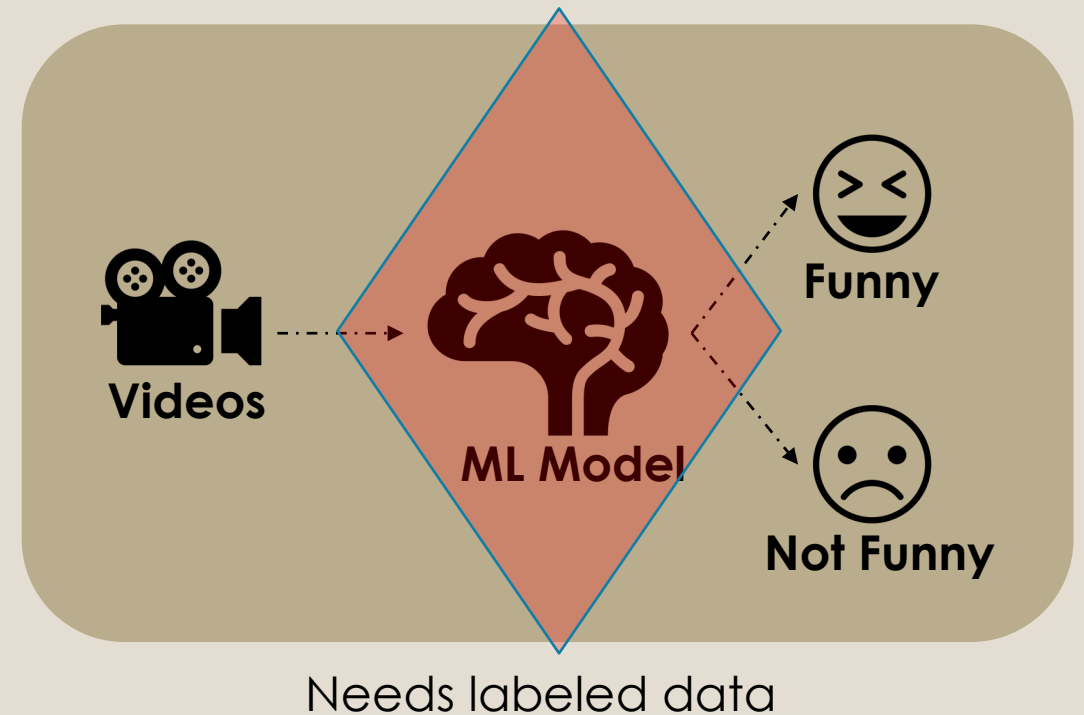
**Instructor** – Dr. Arghyadip Roy

## **Team Members**

Ashish Kumar (224156003)  
Atul Bhagat (224156004)  
Ashish Goswami (224156015)

# What this project is about?

- Goal is to classify movie scenes
- Extracting **affective** information from video
  - Complex Spatial and Temporal dependencies integrated with human perception and information
  - Complex non-linear Process
  - Scene level affective labels convey cumulative information from different modalities
- Efficient search and recommendations



# Why use Reinforcement Learning?


- Difficult to collect a *large annotated corpus*
- Despite the large amount of available movie data, the amount of accurately labeled data is severely limited due to **copyright** and cost of annotation
- Need a learning framework based on RL that is **tolerant to label sparsity**
- Learner should easily make use of any available ground truth in an online fashion

**State Transitions and Rewards Tuple ( $s(t)$ ,  $A(t)$ ,  $s(t + 1)$ ,  $R(t)$ )**

# Challenges with this approach?



Movie Data Complexity




Agent – Environment Interaction

# Challenges with this approach?



**Movie Data Complexity**



**Agent – Environment Interaction**

## Solution

For Simplicity, only **visual modality** is used

And,  $\alpha$  parameter is used, that changes with agent's action

# Data Processing – Extracting Frames

We used **FRIENDS** TV Series to curate a data set for us

**Ffmpeg** – Unix command line tool used to extract frames

Output was rescaled and converted to grayscale

For more information <https://ffmpeg.org/>

```
#!/bin/bash




# Define the input directory and output directory
input_dir="scenes"
output_dir="frames"
wid=256
ht=-1

# Loop over each subdirectory in the input directory
for sub_dir in "$input_dir"/*; do
    # Extract the subdirectory name (i.e., the last component of the path)
    sub_dir_name="$(basename "$sub_dir")"

    # Loop over each scene file in the subdirectory
    for scene_file in "$sub_dir"/*.mkv; do
        # Extract the basename of the scene file (without the extension)
        scene_basename="$(basename "${scene_file%.}*)"

        # Create a subdirectory in the output directory for the frames of this scene
        mkdir -p "$output_dir/$sub_dir_name/$scene_basename"

        # Use Ffmpeg to extract one frame per second from the scene file
        ffmpeg -i "$scene_file" -vf "scale=$wid:$ht,format=gray,fps=20" -r 5 "$output_dir/$sub_dir_name/$scene_basename/frame_%03d.png"
    done
done
```



# Data Processing – Faces Extraction


On the image channel, **most affective information** is contained in the **faces** present in the frame

**Dlib** – Standard Face Detection Library used to extract faces

For multiple faces, one **closest to center** is only considered (assuming, main character is present in the center, thus containing *most relevant affective information*)

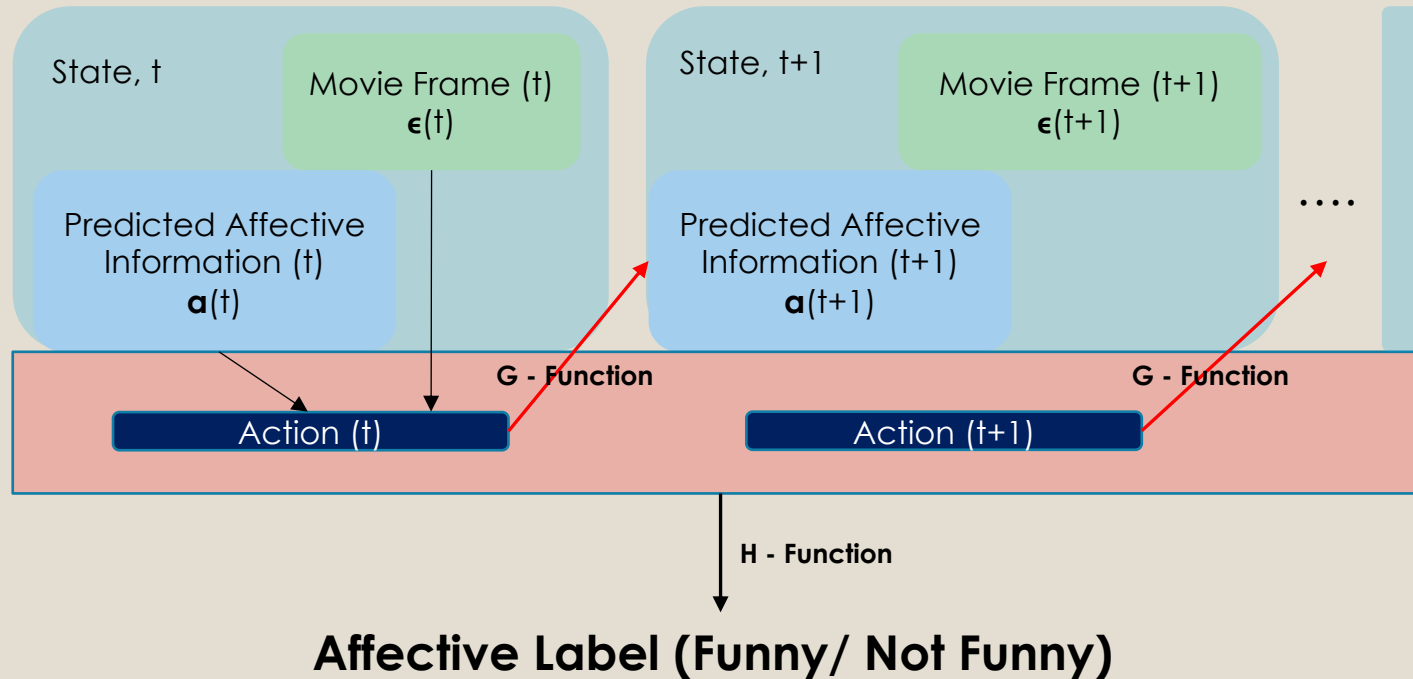


# Facial Expressions Embeddings : CNN





# Coming to the Framework..



$$\mathbf{A}(t) = \mathbf{Q}(s(t)) = \mathbf{Q}([ \epsilon(t), \mathbf{a}(t) ])$$

$$\mathbf{R}(t) = \mathbf{H}( \mathbf{A}(1), \dots, \mathbf{A}(t), \mathbf{F} )$$

$$\mathbf{a}(t + 1) = \mathbf{G}(\mathbf{A}(1), \dots, \mathbf{A}(t))$$

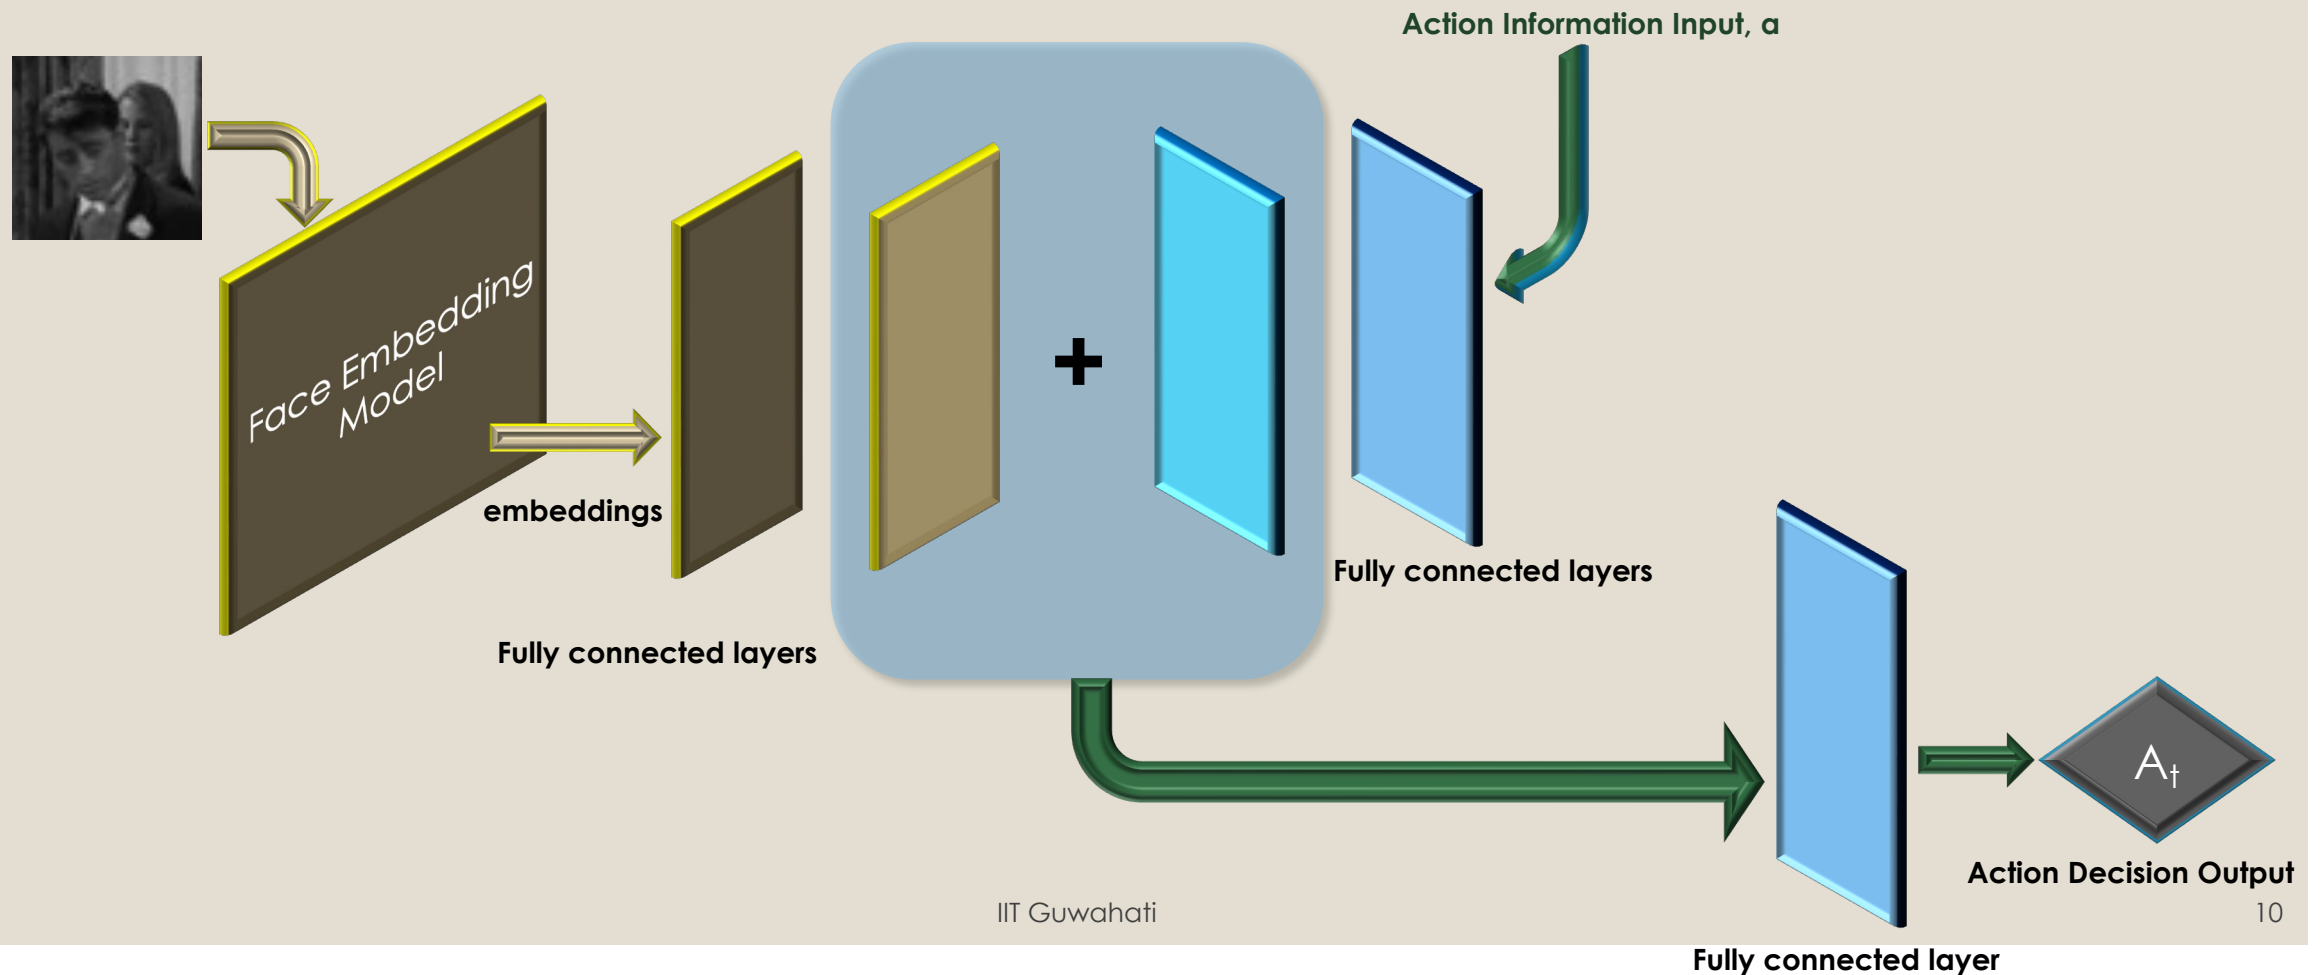
**Q** – Deep Q Network

**F** – True Affective Label



**A**( $t$ ) – Action at timestep,  $t$

Reward Generated by **H** function

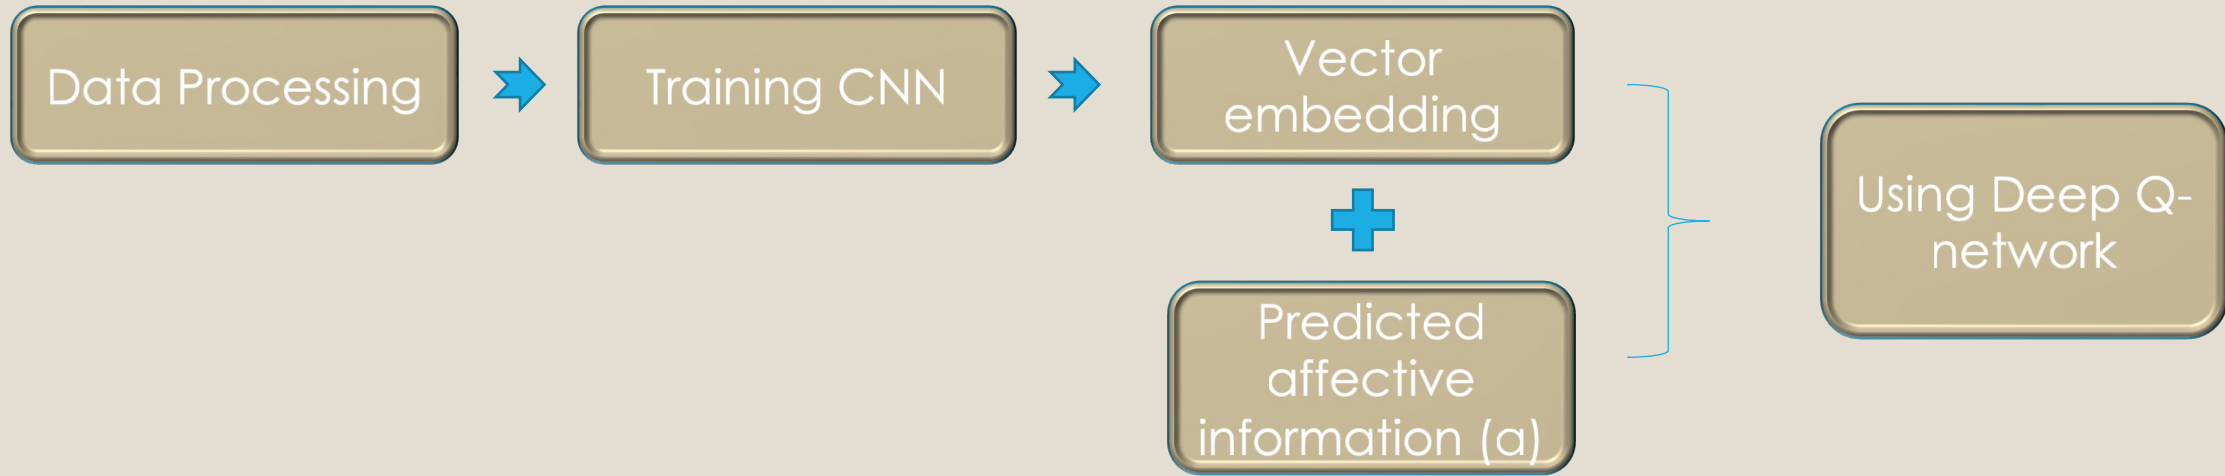
# Framework Continued...



# Q Learning & DQN



$$\text{Loss Function} = [ R_{t+1} + \gamma * \max Q(s_{t+1}, a) ] - Q(s_t, a_t)$$



- 7 classes of emotion
- Frames from video clips using dlib library
- 5 convolution layer (64,128,256)
- 20 epochs
- At each time  $t$ , vector embedding and Predicted affective information is used.
- Agent make decision at time  $t$
- For input  $a(t+1)$  at time  $(t+1)$  either the action output at time  $t$  or human annotations at time  $t$  if available

# Let's have a look at the code : CNN

```
from tensorflow.keras.layers import BatchNormalization
model=Sequential()

model.add(Conv2D(filters=64, kernel_size = (3,3), activation="relu", input_shape=(48,48,1)))
model.add(Conv2D(filters=64, kernel_size = (3,3), activation="relu"))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(BatchNormalization())

model.add(Conv2D(filters=128, kernel_size = (3,3), activation="relu"))
model.add(Conv2D(filters=128, kernel_size = (3,3), activation="relu"))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(BatchNormalization())

model.add(Conv2D(filters=256, kernel_size = (3,3), activation="relu"))
model.add(MaxPooling2D(pool_size=(2,2)))
model.add(BatchNormalization())
model.add(Flatten())
model.add(Dense(512,activation="relu"))

model.add(Dense(7,activation="softmax"))

model.compile(optimizer='adam', loss='categorical_crossentropy',metrics=['accuracy'])

model.summary()
```

**Note:** Only snippets  
are shown here

# Let's have a look at the code : **DLib**

```
import dlib
import cv2

def cropface(frame):
    detector = dlib.get_frontal_face_detector()
    gray = cv2.cvtColor(frame, cv2.COLOR_BGR2GRAY)
    faces = detector(gray, 1)

    if len(faces) == 0:
        return None
    elif len(faces) == 1:
        face = faces[0]
    else:
        center_x, center_y = frame.shape[1] // 2, frame.shape[0] // 2
        dist_to_center = [(f.center().x - center_x) ** 2 + (f.center().y - center_y) ** 2 for f in faces]
        face = faces[dist_to_center.index(min(dist_to_center))]

    left, top, right, bottom = face.left(), face.top(), face.right(), face.bottom()
    return frame[top:bottom, left:right]
```

# Let's have a look at the code : DQN

```
import numpy as np
from keras.layers import Input, Dense, Concatenate

class DQN(tf.keras.Model):
    def __init__(self, action_size, affective_size, learning_rate, name='DQN'):
        super(DQN, self).__init__()
        tf.reset_default_graph()
        self.action_size = action_size
        self.learning_rate = learning_rate
        with tf.compat.v1.variable_scope(name):
            self.actions = tf.keras.Input(shape=(self.action_size), dtype=tf.float32, name='actions')
            self.affective = tf.keras.Input(shape=(affective_size), dtype=tf.float32, name='affective')
            self.affective_fc1 = tf.layers.Dense(units=64, activation='relu')
            self.output1 = tf.layers.Dense(units=self.action_size, activation='linear')
            self.affective_fc2 = tf.layers.Dense(units=32, activation='relu')
            self.fused_fc1 = tf.layers.Dense(units=64, activation='relu')
            self.fused_fc2 = tf.layers.Dense(units=32, activation='relu')
        def call(self, embedding, action):
            action = np.reshape(action, (1, 2))
            # print(actions)
            x = self.affective_fc1(action)
            x = self.affective_fc2(x)
            x = Concatenate(axis=-1)([x, embedding])
            x = self.fused_fc1(x)
            x = self.fused_fc2(x)
            # Q-value output layer
            x = self.output1(x)
            # Q-value for selected action
            Q = tf.reduce_sum(tf.multiply(x, action), axis=1)
            return Q

        def predict(self, state, affective):
            return self.session.run(self.output, feed_dict={self.inputs: state, self.affective: affective})
```

```
class DQNAgent():

    def __init__(self, action_size=2,
                 discount=0.99,
                 eps_max=1.0,
                 eps_min=0.01,
                 eps_decay=0.995,
                 memory_capacity=5000,
                 lr=1e-3,
                 train_mode=True):

        # for epsilon-greedy exploration strategy
        self.epsilon = eps_max
        self.epsilon_min = eps_min
        self.epsilon_decay = eps_decay
        self.affective_size=2

        # for defining how far-sighted or myopic the agent should be
        self.discount = discount

        self.action_size = action_size

        # instances of the network for current policy and its target
        opt = tf.keras.optimizers.Adam(learning_rate=lr)
        self.policy_net = DQN(self.action_size, self.affective_size, lr)
        self.target_net = DQN(self.action_size, self.affective_size, lr)
        self.policy_net.compile(loss='mse', optimizer=opt)
        self.target_net.compile(loss='mse', optimizer=opt)
        # self.target_net.eval()
        self.target_net.set_weights(self.policy_net.get_weights()) # since no
        if not train_mode:
            self.policy_net.trainable = False
        # instance of the replay buffer
        self.memory = ReplayMemory(capacity=memory_capacity)
```

# Let's have a look at the code : DQN

```
def learn(self, batchsize, embedding, actions1, model, next_state):  
  
    # select n samples picked uniformly at random from the experience replay memory, such that n=batchsize  
    if len(self.memory) < batchsize:  
        return  
  
    # print(memory.pop())  
    # state, action, rewards, next_states, dones = self.memory.sample(batchsize)  
    # get q values of the actions that were taken, i.e calculate qpred;  
    # actions vector has to be explicitly reshaped to nx1-vector  
    state1=np.reshape(next_state,(48,48,3))  
    img= cv2.cvtColor(state1,cv2.COLOR_BGR2GRAY)  
    img = img.astype("float32")  
    img=img/255.0  
    embd=model.predict(np.asarray([img]))  
    actions=np.zeros(2)  
    action=self.select_action(next_state)  
    actions[action]=1.0  
    next_state, reward, done, info = env.step(actions)  
    with tf.GradientTape() as tape:  
        q_pred = self.policy_net.call(embedding,actions1)  
        # calculate target q-values, such that  $y_j = r_j + q(s', a')$ , but if current state is a terminal state, then  $y_j = r_j$   
        q_target =self.target_net.call(embd,actions)  
        q_target = tf.where(done, tf.zeros_like(q_target), q_target)  
        y_j = reward + (self.discount * q_target)  
        y_j = tf.reshape(y_j, (-1, 1))  
        # calculate the loss as the mean-squared error of yj and qpred  
  
        loss = tf.keras.losses.MSE(y_j, q_pred)  
        print("loss=" ,loss)  
    gradients = tape.gradient(loss, self.policy_net.trainable_variables)  
    self.policy_net.optimizer.apply_gradients(zip(gradients, self.policy_net.trainable_variables))
```



# Let's have a look at the code : DQN

```
def train(env,dqn_agent,dataset,actions,num_train_eps,update_frequency,batchsize,results_basepath):
    # print(actions,dataset[0].shape)
    fill_memory(env, dqn_agent, dataset,actions)
    print('Memory filled. Current capacity: ', len(dqn_agent.memory))

    reward_history = []
    epsilon_history = []
    step_cnt = 0
    best_score = -np.inf

    for ep_cnt in range(num_train_eps):
        epsilon_history.append(dqn_agent.epsilon)
        # print(dqn_agent.memory.sample(1)[3])
        for i in range(len(dataset)):
            done = False
            state = env.reset(dataset[i])
            state1=np.reshape(state,(48,48,3))
            img= cv2.cvtColor(state1,cv2.COLOR_BGR2GRAY)
            img = img.astype("float32")
            img=img/255.0
            embd=embedding_model.predict(np.asarray([img]))
            ep_score = 0

            while not done:
                print(ep_cnt,i)
                action = dqn_agent.select_action(img)
                actions=np.zeros(2)
                actions[action]=1.0
                actions=actions.astype("float32")
                next_state, reward, done, info = env.step(actions)
                dqn_agent.memory.push(state=state, action=actions, next_state=next_state, reward=reward, done=done)
                dqn_agent.learn(batchsize=batchsize,embedding=embd,actions1=actions,model=embedding_model,next_state=next_state)
```

# Let's have a look at the code : **DQN**

```
if step_cnt % update_frequency == 0:
    dqn_agent.update_target_net()

state = next_state
ep_score += reward
step_cnt += 1

dqn_agent.update_epsilon()

reward_history.append(ep_score)
current_avg_score = np.mean(reward_history[-100:]) # moving average of last 100 episodes

print('Ep: {}, Total Steps: {}, Ep: Score: {}, Avg score: {}; Epsilon: {}'.format(ep_cnt, step_cnt, ep_score, current_avg_score, epsilon_history[-1]))

if current_avg_score >= best_score:
    dqn_agent.save_model('{}dqn_model'.format(results_basepath))
    best_score = current_avg_score

with open('{}train_reward_history.pkl'.format(results_basepath), 'wb') as f:
    pickle.dump(reward_history, f)

with open('{}train_epsilon_history.pkl'.format(results_basepath), 'wb') as f:
    pickle.dump(epsilon_history, f)
```

# Let's have a look at the code : DQN

```
def train(env,dqn_agent,dataset,actions,num_train_eps,update_frequency,batchsize,results_basepath):
    # print(actions,dataset[0].shape)
    fill_memory(env, dqn_agent, dataset,actions)
    print('Memory filled. Current capacity: ', len(dqn_agent.memory))

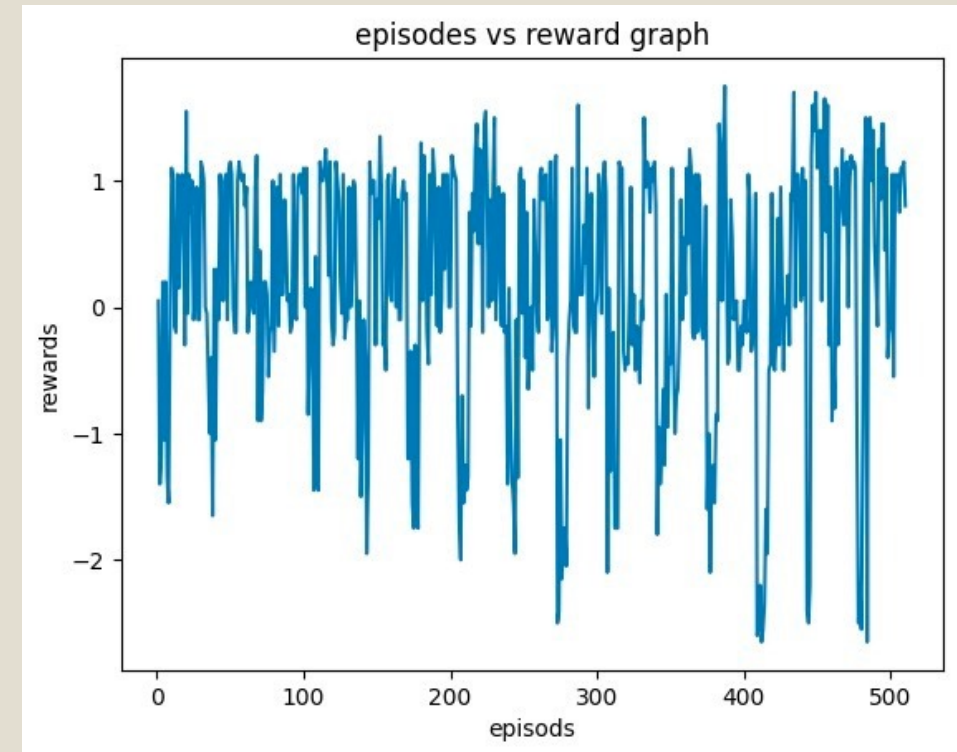
    reward_history = []
    epsilon_history = []
    step_cnt = 0
    best_score = -np.inf

    for ep_cnt in range(num_train_eps):
        epsilon_history.append(dqn_agent.epsilon)
        # print(dqn_agent.memory.sample(1)[3])
        for i in range(len(dataset)):
            done = False
            state = env.reset(dataset[i])
            state1=np.reshape(state,(48,48,3))
            img= cv2.cvtColor(state1,cv2.COLOR_BGR2GRAY)
            img = img.astype("float32")
            img=img/255.0
            embd=embedding_model.predict(np.asarray([img]))
            ep_score = 0

            while not done:
                print(ep_cnt,i)
                action = dqn_agent.select_action(img)
                actions=np.zeros(2)
                actions[action]=1.0
                actions=actions.astype("float32")
                next_state, reward, done, info = env.step(actions)
                dqn_agent.memory.push(state=state, action=actions, next_state=next_state, reward=reward, done=done)
                dqn_agent.learn(batchsize=batchsize,embedding=embd,actions1=actions,model=embedding_model,next_state=next_state)
```

# Model Training and Results...

```
1/1 [=====] - 0s 31ms/step
actions= tf.Tensor([1.0866777], shape=(1,), dtype=float32) tf.Tensor([[1.0866777 0.94620126]], shape=(1, 2), dtype=float32)
1/1 [=====] - 0s 33ms/step
actions= tf.Tensor([0.8867666], shape=(1,), dtype=float32) tf.Tensor([[0.8867666 0.7592584]], shape=(1, 2), dtype=float32)
loss= tf.Tensor([0.00751302], shape=(1,), dtype=float32)
Ep: 14, Total Steps: 7956, Ep: Score: 1.1, Avg score: 0.16949999999999996; Epsilon: 0.09199970504166631
14 32
1/1 [=====] - 0s 27ms/step
actions= tf.Tensor([0.9435819], shape=(1,), dtype=float32) tf.Tensor([[0.9435819 0.8060961]], shape=(1, 2), dtype=float32)
1/1 [=====] - 0s 28ms/step
actions= tf.Tensor([0.9435819], shape=(1,), dtype=float32) tf.Tensor([[0.9435819 0.8060961]], shape=(1, 2), dtype=float32)
loss= tf.Tensor([0.00115699], shape=(1,), dtype=float32)
14 32
1/1 [=====] - 0s 29ms/step
actions= tf.Tensor([0.94522446], shape=(1,), dtype=float32) tf.Tensor([[0.94522446 0.80432075]], shape=(1, 2), dtype=float32)
1/1 [=====] - 0s 27ms/step
actions= tf.Tensor([0.93341285], shape=(1,), dtype=float32) tf.Tensor([[0.93341285 0.7978898 ]], shape=(1, 2), dtype=float32)
loss= tf.Tensor([0.0030528], shape=(1,), dtype=float32)
14 32
1/1 [=====] - 0s 29ms/step
actions= tf.Tensor([0.93278104], shape=(1,), dtype=float32) tf.Tensor([[0.93278104 0.7949435 ]], shape=(1, 2), dtype=float32)
1/1 [=====] - 0s 30ms/step
actions= tf.Tensor([0.92538965], shape=(1,), dtype=float32) tf.Tensor([[0.92538965 0.7878702 ]], shape=(1, 2), dtype=float32)
loss= tf.Tensor([0.00451839], shape=(1,), dtype=float32)
14 32
1/1 [=====] - 0s 30ms/step
actions= tf.Tensor([0.92754316], shape=(1,), dtype=float32) tf.Tensor([[0.92754316 0.7867954 ]], shape=(1, 2), dtype=float32)
1/1 [=====] - 0s 27ms/step
actions= tf.Tensor([0.92754316], shape=(1,), dtype=float32) tf.Tensor([[0.92754316 0.7867954 ]], shape=(1, 2), dtype=float32)
loss= tf.Tensor([0.00524999], shape=(1,), dtype=float32)
Ep: 14, Total Steps: 7960, Ep: Score: 1.15, Avg score: 0.20699999999999996; Epsilon: 0.09199970504166631
14 33
```



# How to improve on this?

- Use **3D CNN**, that could capture the temporal dependency in a better fashion
- **Create a DataSet** from old movies/ TV Shows, out of the copyright
- Create a **online task** for training, like training on YouTube shorts, Instagram reels and extracting labels from the comments
- Include **audio modality** in the model
- Use a **knowledge base** to understand sarcasm/ puns etc. from subtitles, as well as audio



# THANK YOU

***Torture the data, and it will confess to anything ~ Ronald Coase***